

# N-grams : Tidytext Vs tm

This notebook demonstrates a difference I found in computing n-grams from the 'tidytext' package versus the 'tm' package, when the corpus contains multiple documents/lines.

Create a little example corpus, where EOS signifies end of sentence:

```
rm(list=ls())
corpus <- c("buy the book EOS", "buy the book EOS")
corpus
```

```
## [1] "buy the book EOS" "buy the book EOS"
```

Tokenize w/ tidytext:

```
library(dplyr)
```

```
##
## Attaching package: 'dplyr'
## The following objects are masked from 'package:stats':
##
##   filter, lag
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union
```

```
library(tidytext)
```

```
df <- data_frame(corpus)
uni <- unnest_tokens(df, ngram, corpus, token="words")
bi <- unnest_tokens(df, ngram, corpus, token="ngrams", n=2)
tri <- unnest_tokens(df, ngram, corpus, token="ngrams", n=3)
```

Tokenize w/ tm:

```
library(tm)
```

```
## Loading required package: NLP
```

```
library(RWeka)
```

```
mysource <- VectorSource(corpus)
corp <- VCorpus(mysource)
```

```
# find bigrams
```

```
bigramTokenizer <- function(x) unlist(lapply(ngrams(words(x), 2), paste, collapse = " "), use.names = F))
bi_tm <- TermDocumentMatrix(corp, control = list(tokenize = bigramTokenizer))
```

```
# find trigrams
```

```
trigramTokenizer <- function(x) unlist(lapply(ngrams(words(x), 3), paste, collapse = " "), use.names = F))
tri_tm <- TermDocumentMatrix(corp, control = list(tokenize = trigramTokenizer))
```

We see that there are 4 unique bigrams from the tidytext method:

```
bi_ct <- bi %>% count(ngram, sort=TRUE)
bi_ct
```

```
## # A tibble: 4 × 2
```

```
##      ngram      n
##      <chr> <int>
## 1 book eos      2
## 2 buy the      2
## 3 the book     2
## 4 eos buy      1
```

While there are only 3 unique bigrams from the tm package:

```
data_frame(rownames(as.matrix(bi_tm)))
```

```
## # A tibble: 3 × 1
##   `rownames(as.matrix(bi_tm))`
##   <chr>
## 1 book eos
## 2 buy the
## 3 the book
```

What I think happened is that the tm package computed n-grams for each ‘document’ or line of the corpus, while the tidytext package computed n-grams as if all the documents were one long line (the bigram ‘eos buy’ spans the 2 lines of the corpus). If we were to remove the tidytext n-grams where EOS is at the beginning or middle of a sentence, the results would match.

```
sessionInfo()
```

```
## R version 3.3.2 (2016-10-31)
## Platform: x86_64-apple-darwin13.4.0 (64-bit)
## Running under: OS X Yosemite 10.10.5
##
## locale:
## [1] en_US.UTF-8/en_US.UTF-8/en_US.UTF-8/C/en_US.UTF-8/en_US.UTF-8
##
## attached base packages:
## [1] stats      graphics  grDevices  utils      datasets  methods   base
##
## other attached packages:
## [1] RWeka_0.4-29  tm_0.6-2      NLP_0.1-9      tidytext_0.1.2
## [5] dplyr_0.5.0
##
## loaded via a namespace (and not attached):
## [1] Rcpp_0.12.7      knitr_1.15      magrittr_1.5
## [4] RWekajars_3.9.0-1 mnormt_1.5-5    lattice_0.20-34
## [7] R6_2.2.0         stringr_1.1.0   plyr_1.8.4
## [10] tools_3.3.2     parallel_3.3.2  grid_3.3.2
## [13] broom_0.4.1     nlme_3.1-128    psych_1.6.9
## [16] DBI_0.5-1       janeaustenr_0.1.4 htmltools_0.3.5
## [19] lazyeval_0.2.0  yaml_2.1.13     assertthat_0.1
## [22] digest_0.6.10  tibble_1.2      Matrix_1.2-7.1
## [25] rJava_0.9-8     purrr_0.2.2     tidyr_0.6.0
## [28] reshape2_1.4.2 tokenizers_0.1.4 SnowballC_0.5.1
## [31] slam_0.1-38     evaluate_0.10   rmarkdown_1.1
## [34] stringi_1.1.2   foreign_0.8-67
```